

Problem Sheet 1

You should try to your best to do the core problems. If time permits, try and do the rest as well.

Problem 1. Arrow matrix-vector multiplication (core problem)

Consider the multiplication of the two “arrow matrices” A with a vector x , implemented as a function `arrowmatvec(d, a, x)` in the following MATLAB script

Listing 1: multiplying a vector with the product of two “arrow matrices”

```
1 function y = arrowmatvec(d, a, x)
2 % Multiplying a vector with the product of two ``arrow
   matrices''
3 % Arrow matrix is specified by passing two column
   vectors a and d
4 if (length(d) ~= length(a)), error('size mismatch'); end
5 % Build arrow matrix using the MATLAB function diag()
6 A = [diag(d(1:end-1)), a(1:end-1); (a(1:end-1))', d(end)];
7 y = A*A*x;
```

(1a) ☐ For general vectors $d = (d_1, \dots, d_n)^\top$ and $a = (a_1, \dots, a_n)^\top$, sketch the matrix A created in line 6 of Listing 1.

HINT: This MATLAB script is provided as file `arrowmatvec.m`.

(1b) ☐ The `tic-toc` timing results for `arrowmatvec.m` are available in Figure 1. Give a detailed explanation of the results.

HINT: This MATLAB created figure is provided as file `arrowmatvectiming.{eps, jpg}`.

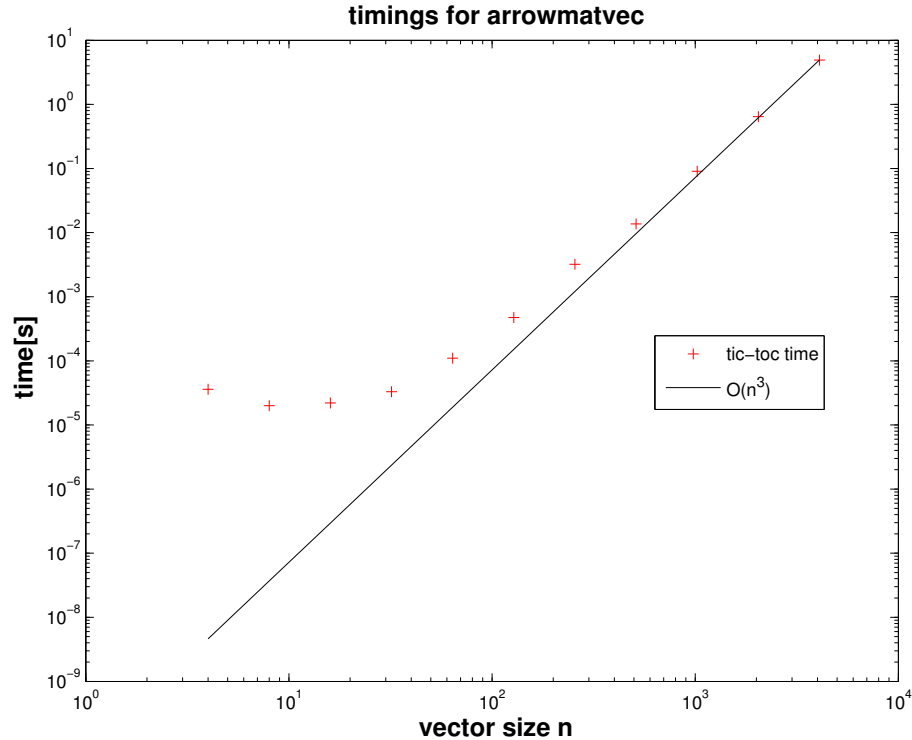


Figure 1: timings for `arrowmatvec(d, a, x)`

(1c) ☐ Write an *efficient* MATLAB function

```
function y = arrowmatvec2(d, a, x)
```

that computes the same multiplication as in code 1 but with optimal asymptotic complexity with respect to n . Here \mathbf{d} passes the vector $(d_1, \dots, d_n)^T$ and \mathbf{a} passes the vector $(a_1, \dots, a_n)^T$.

(1d) ☐ What is the complexity of your algorithm from sub-problem (1c) (with respect to problem size n)?

(1e) ☐ Compare the runtime of your implementation and the implementation given in code 1 for $n = 2^5, 6, \dots, 12$. Use the routines `tic` and `toc` as explained in example [1, Ex. 1.4.10] of the Lecture Slides.

(1f) ☐ Write the EIGEN codes corresponding to the functions `arrowmatvec` and `arrowmatvec2`.

Problem 2. Avoiding cancellation (core problem)

In [1, Section 1.5.4] we saw that the so-called *cancellation phenomenon* is a major cause of numerical instability, cf. [1, § 1.5.38]. Cancellation is the massive amplification of *relative errors* when subtracting two real numbers of about the same value.

Fortunately, expressions vulnerable to cancellation can often be recast in a mathematically equivalent form that is no longer affected by cancellation, see [1, § 1.5.43]. There we studied several examples, and this problem gives some more,

(2a) ◻ We consider the function

$$f_1(x_0, h) := \sin(x_0 + h) - \sin(x_0) . \quad (1)$$

It can be transformed into another form, $f_2(x_0, h)$, using the trigonometric identity

$$\sin(\varphi) - \sin(\psi) = 2 \cos\left(\frac{\varphi + \psi}{2}\right) \sin\left(\frac{\varphi - \psi}{2}\right) .$$

Thus, f_1 and f_2 give the same values, in exact arithmetic, for any given argument values x_0 and h .

1. Derive $f_2(x_0, h)$, which does no longer involve the difference of return values of trigonometric functions.
2. Suggest a formula that avoids cancellation errors for computing the approximation $(f(x_0 + h) - f(x_0))/h$ of the derivative of $f(x) := \sin(x)$ at $x = x_0$. Write a MATLAB program that implements your formula and computes an approximation of $f'(1.2)$, for $h = 1 \cdot 10^{-20}, 1 \cdot 10^{-19}, \dots, 1$.
HINT: For background information refer to [1, Ex. 1.5.40].
3. Plot the error (in doubly logarithmic scale using MATLAB's `loglog` plotting function) of the derivative computed with the suggested formula and with the naive implementation using f_1 .
4. Explain the observed behaviour of the error.

(2b) ◻ Using a trick applied in [1, Ex. 1.5.48] show that

$$\ln(x - \sqrt{x^2 - 1}) = -\ln(x + \sqrt{x^2 - 1}) .$$

Which of the two formulas is more suitable for numerical computation? Explain why, and provide a numerical example in which the difference in accuracy is evident.

(2c) ◻ For the following expressions, state the numerical difficulties that may occur, and rewrite the formulas in a way that is more suitable for numerical computation.

1. $\sqrt{x + \frac{1}{x}} - \sqrt{x - \frac{1}{x}}$, where $x \gg 1$.

2. $\sqrt{\frac{1}{a^2} + \frac{1}{b^2}}$, where $a \approx 0, b \approx 1$.

Problem 3. Kronecker product

In [1, Def. 1.4.16] we learned about the so-called Kronecker product, available in MATLAB through the command `kron`. In this problem we revisit the discussion of [1, Ex. 1.4.17]. Please refresh yourself on this example and study [1, Code 1.4.18] again.

As in [1, Ex. 1.4.17], the starting point is the line of MATLAB code

$$\mathbf{y} = \text{kron}(\mathbf{A}, \mathbf{B}) * \mathbf{x}, \quad (2)$$

where the arguments are $\mathbf{A}, \mathbf{B} \in \mathbb{R}^{n,n}, \mathbf{x} \in \mathbb{R}^{n \cdot n}$.

(3a) ◻ Obtain further information about the `kron` command from MATLAB help issuing `doc kron` in the MATLAB command window.

(3b) ◻ Explicitly write Eq. (2) in the form $\mathbf{y} = \mathbf{M}\mathbf{x}$ (i.e. write down \mathbf{M}), for $\mathbf{A} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$ and $\mathbf{B} = \begin{pmatrix} 5 & 6 \\ 7 & 8 \end{pmatrix}$.

(3c) ◻ What is the asymptotic complexity (\rightarrow [1, Def. 1.4.3]) of the MATLAB code (2)? Use the Landau symbol from [1, Def. 1.4.4] to state your answer.

(3d) ◻ Measure the runtime of (2) for $n = 2^{3,4,5,6}$ and random matrices. Use the MATLAB functions `tic` and `toc` as explained in example [1, Ex. 1.4.10] of the Lecture Slides.

(3e) ◻ Explain in detail, why (2) can be replaced with the single line of MATLAB code

$$\mathbf{y} = \text{reshape}(\mathbf{B} * \text{reshape}(\mathbf{x}, n, n) * \mathbf{A}', n * n, 1); \quad (3)$$

and compare the execution times of (2) and (3) for random matrices of size $n = 2^{3,4,5,6}$.

(3f) ◻ Based on the EIGEN numerical library (\rightarrow [1, Section 1.2.2]) implement a C++ function

```

template <class Matrix>
void kron(const Matrix & A, const Matrix & B, Matrix &
    C) {
    // Your code here
}

```

returns the Kronecker product of the argument matrices A and B in the matrix C.

HINT: Feel free (but not forced) to use the partial codes provided in `kron.cpp` as well as the CMake file `CMakeLists.txt` (including `cmake-modules`) and the timing header file `timer.h`.

(3g) ☐ Devise an implementation of the MATLAB code (2) in C++ according to the function definition

```

template <class Matrix, class Vector>
void kron_mv(const Matrix & A, const Matrix & B, const
    Vector & x, Vector & y);

```

The meaning of the arguments should be self-explanatory.

(3h) ☐ Now, using a function definition similar to that of the previous sub-problem, implement the C++ equivalent of (3) in the function `kron_mv_fast`.

HINT: Study [1, Rem. 1.2.18] about “reshaping” matrices in EIGEN.

(3i) ☐ Compare the runtimes of your two implementations as you did for the MATLAB implementations in sub-problem (3e).

Problem 4. Structured matrix–vector product

In [1, Ex. 1.4.14] we saw how the particular structure of a matrix can be exploited to compute a matrix-vector product with substantially reduced computational effort. This problem presents a similar case.

Consider the real $n \times n$ matrix \mathbf{A} defined by $(\mathbf{A})_{i,j} = a_{i,j} = \min\{i, j\}$, for $i, j = 1, \dots, n$. The matrix-vector product $\mathbf{y} = \mathbf{A}\mathbf{x}$ can be implemented in MATLAB as

$$\mathbf{y} = \min(\text{ones}(n, 1) * (1:n), (1:n)' * \text{ones}(1, n)) * \mathbf{x}; \quad (4)$$


(4a) ☐ What is the asymptotic complexity (for $n \rightarrow \infty$) of the evaluation of the MATLAB command displayed above, with respect to the problem size parameter n ?


(4b)  Write an *efficient* MATLAB function


```
function y = multAmin(x)
```

that computes the same multiplication as (4) but with a better asymptotic complexity with respect to n .

HINT: you can test your implementation by comparing the returned values with the ones obtained with code (4).


(4c)  What is the asymptotic complexity (in terms of problem size parameter n) of your function `multAmin`?

(4d)  Compare the runtime of your implementation and the implementation given in (4) for $n = 2^5, 6, \dots, 12$. Use the routines `tic` and `toc` as explained in example [1, Ex. 1.4.10] of the Lecture Slides.

(4e)  Can you solve task (4b) without using any `for`- or `while`-loop? Implement it in the function

```
function y = multAmin2(x)
```

HINT: you may use the MATLAB built-in command `cumsum`.


(4f)  Consider the following MATLAB script `multAB.m`:

Listing 2: MATLAB script calling `multAmin`


```
1 n = 10;
2 B = diag(-ones(n-1,1),-1)+diag([2*ones(n-1,1);1],0)...
3     + diag(-ones(n-1,1),1);
4 x = rand(n,1);
5 fprintf('||x-y|| = %d\n', norm(multAmin(B*x)-x));
```

Sketch the matrix **B** created in line 3 of `multAB.m`.

HINT: this MATLAB script is provided as file `multAB.m`.

(4g)  Run the code of Listing 2 several times and conjecture a relationship between the matrices **A** and **B** from the output. Prove your conjecture.


HINT: You must take into account that computers inevitably commit round-off errors, see [1, Section 1.5].

(4h)  Implement a C++ function with declaration

```
1 template <class Vector>
2 void minmatmv(const Vector &x, Vector &y);
```

that realizes the efficient version of the MATLAB line of code (4). Test your function by comparing with output from the equivalent MATLAB functions.

Problem 5. Matrix powers


(5a)  Implement a MATLAB function


Pow(A, k)

that, using only basic linear algebra operations (including matrix-vector or matrix-matrix multiplications), computes efficiently the k^{th} power of the $n \times n$ matrix **A**.

HINT: use the MATLAB operator \wedge to test your implementation on random matrices **A**.

HINT: use the MATLAB functions `de2bi` to extract the “binary digits” of an integer.


(5b)  Find the asymptotic complexity in k (and n) taking into account that in MATLAB a matrix-matrix multiplication requires a $O(n^3)$ effort.

(5c)  Plot the runtime of the built-in MATLAB power (\wedge) function and find out the complexity. Compare it with the function `Pow` from (5a).

Use the matrix

$$A_{j,k} = \frac{1}{\sqrt{n}} \exp\left(\frac{2\pi i j k}{n}\right)$$

to test the two functions.

(5d)  Using EIGEN, devise a C++ function with the calling sequence

```
1 template <class Matrix>
2 void matPow(const Matrix &A, unsigned int k);
```

that computes the k^{th} power of the square matrix **A** (passed in the argument **A**). Of course, your implementation should be as efficient as the MATLAB version from sub-problem (5a).

HINT: matrix multiplication suffers no aliasing issues (you can safely write $A = A \star A$).

HINT: feel free to use the provided `matPow.cpp`.

HINT: you may want to use `log` and `ceil`.

HINT: EIGEN implementation of power (`A.pow(k)`) can be found in:

```
#include <unsupported/Eigen/MatrixFunctions>
```

Problem 6. Complexity of a MATLAB function

In this problem we recall a concept from linear algebra, the diagonalization of a square matrix. Unless you can still define what this means, please look up the chapter on “eigenvalues” in your linear algebra lecture notes. This problem also has a subtle relationship with Problem 5.

We consider the MATLAB function defined in `getit.m` (cf. Listing 3)

Listing 3: MATLAB implementation of `getit` for Problem 6..

```
1 function y = getit(A, x, k)
2     [S,D] = eig(A);
3     y = S*diag(diag(D).^k)*(S\'x);
4 end
```

HINT: Give the command `doc eig` in MATLAB to understand what `eig` does.

HINT: You may use that `eig` applied to an $n \times n$ -matrix requires an asymptotic computational effort of $O(n^3)$ for $n \rightarrow \infty$.

HINT: in MATLAB, the function `diag(x)` for $x \in \mathbb{R}^n$, builds a diagonal, $n \times n$ matrix with x as diagonal. If M is a $n \times n$ matrix, `diag(M)` returns (extracts) the diagonal of M as a vector in \mathbb{R}^n .

HINT: the operator $v.^k$ for $v \in \mathbb{R}^n$ and $k \in \mathbb{N} \setminus \{0\}$ returns the vector with components v_i^k (i.e. component-wise exponent)

(6a) ☐ What is the output of `getit`, when A is a diagonalizable $n \times n$ matrix, $x \in \mathbb{R}^n$ and $k \in \mathbb{N}$?

(6b) ☐ Fix $k \in \mathbb{N}$. Discuss (in detail) the asymptotic complexity of `getit` $n \rightarrow \infty$.

Issue date: 17.09.2015

Hand-in: 24.09.2015 (in the boxes in front of HG G 53/54).

Version compiled on: September 19, 2015 (v. 1.1).

References

- [1] R. Hiptmair. *Lecture slides for course "Numerical Methods for CSE"*.
<http://www.sam.math.ethz.ch/~hiptmair/tmp/NumCSE/NumCSE15.pdf>. 2015.